

convnets_mnist_dataset

February 1, 2024

```
[ ]: import keras
keras.__version__
```

```
[ ]: '2.15.0'
```

We will use our convnet to classify MNIST digits, a task that you've already been through in Chapter 2, using a densely-connected network (our test accuracy then was 97.8%).

The 6 lines of code below show you what a basic convnet looks like. It's a stack of `Conv2D` and `MaxPooling2D` layers. We'll see in a minute what they do concretely. Importantly, a convnet takes as input tensors of shape `(image_height, image_width, image_channels)` (not including the batch dimension). In our case, we will configure our convnet to process inputs of size `(28, 28, 1)`, which is the format of MNIST images. We do this via passing the argument `input_shape=(28, 28, 1)` to our first layer.

```
[1]: from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

Let's display the architecture of our convnet so far:

```
[2]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|------------------------------|--------------------|---------|
| conv2d (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 11, 11, 64) | 18496 |

```
max_pooling2d_1 (MaxPoolin (None, 5, 5, 64)          0
g2D)

conv2d_2 (Conv2D)          (None, 3, 3, 64)          36928
```

```
=====
Total params: 55744 (217.75 KB)
Trainable params: 55744 (217.75 KB)
Non-trainable params: 0 (0.00 Byte)
-----
```

You can see above that the output of every `Conv2D` and `MaxPooling2D` layer is a 3D tensor of shape (height, width, channels). The width and height dimensions tend to shrink as we go deeper in the network. The number of channels is controlled by the first argument passed to the `Conv2D` layers (e.g. 32 or 64).

The next step would be to feed our last output tensor (of shape (3, 3, 64)) into a densely-connected classifier network like those you are already familiar with: a stack of `Dense` layers. These classifiers process vectors, which are 1D, whereas our current output is a 3D tensor. So first, we will have to flatten our 3D outputs to 1D, and then add a few `Dense` layers on top:

```
[3]: model.add(layers.Flatten())
      model.add(layers.Dense(64, activation='relu'))
      model.add(layers.Dense(10, activation='softmax'))
```

We are going to do 10-way classification, so we use a final layer with 10 outputs and a softmax activation. Now here's what our network looks like:

```
[4]: model.summary()
```

```
Model: "sequential"
```

```
-----
Layer (type)                Output Shape          Param #
-----
conv2d (Conv2D)             (None, 26, 26, 32)    320
max_pooling2d (MaxPooling2D) (None, 13, 13, 32)    0
conv2d_1 (Conv2D)           (None, 11, 11, 64)    18496
max_pooling2d_1 (MaxPoolin (None, 5, 5, 64)      0
g2D)
conv2d_2 (Conv2D)           (None, 3, 3, 64)      36928
flatten (Flatten)           (None, 576)           0
```

| | | |
|-----------------|------------|-------|
| dense (Dense) | (None, 64) | 36928 |
| dense_1 (Dense) | (None, 10) | 650 |

```
=====
Total params: 93322 (364.54 KB)
Trainable params: 93322 (364.54 KB)
Non-trainable params: 0 (0.00 Byte)
-----
```

As you can see, our (3, 3, 64) outputs were flattened into vectors of shape (576,), before going through two Dense layers.

Now, let's train our convnet on the MNIST digits. We will reuse a lot of the code we have already covered in the MNIST example from Chapter 2.

```
[5]: from keras.datasets import mnist
      from keras.utils import to_categorical
      from sklearn.model_selection import train_test_split

      (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

      # Reshape the images for compatibility with Conv2D layer
      #train_images = train_images.reshape((train_images.shape[0], 28, 28, 1))

      train_images = train_images.reshape((60000, 28, 28, 1))
      train_images = train_images.astype('float32') / 255

      test_images = test_images.reshape((10000, 28, 28, 1))
      test_images = test_images.astype('float32') / 255

      train_labels = to_categorical(train_labels,num_classes=10)
      test_labels = to_categorical(test_labels,num_classes=10)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
```

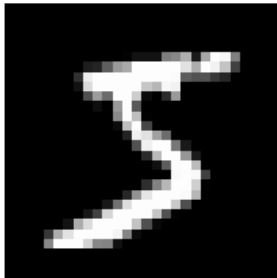
```
[6]: print(train_labels.shape)
      print(test_labels.shape)
```

```
(60000, 10)
(10000, 10)
```

```
[7]: import matplotlib.pyplot as plt
# Display the first four images
plt.figure(figsize=(10, 5))
for i in range(4):
    plt.subplot(2, 2, i + 1)
    plt.imshow(train_images[i], cmap='gray')
    plt.title(f"Label: {train_labels[i]}")
    plt.axis('off')

plt.show()
```

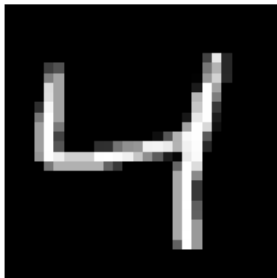
Label: [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]



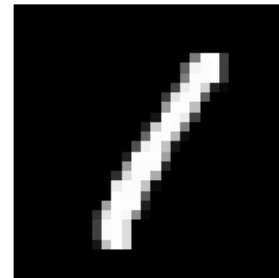
Label: [1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]



Label: [0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]



Label: [0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]



```
[9]: model.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
history = model.fit(train_images, train_labels, epochs=5, batch_size=64) #,
↪ validation_data=(val_images, val_labels))
```

Epoch 1/5

938/938 [=====] - 49s 51ms/step - loss: 0.1838 - accuracy: 0.9417

Epoch 2/5

938/938 [=====] - 48s 51ms/step - loss: 0.0470 - accuracy: 0.9853

Epoch 3/5

938/938 [=====] - 49s 52ms/step - loss: 0.0322 - accuracy: 0.9898

```
Epoch 4/5
938/938 [=====] - 47s 50ms/step - loss: 0.0238 -
accuracy: 0.9927
Epoch 5/5
938/938 [=====] - 48s 51ms/step - loss: 0.0185 -
accuracy: 0.9941
```

Let's evaluate the model on the test data:

```
[10]: test_loss, test_acc = model.evaluate(test_images, test_labels)
```

```
313/313 [=====] - 3s 9ms/step - loss: 0.0268 -
accuracy: 0.9931
```

```
[11]: test_acc
```

```
[11]: 0.9930999875068665
```